

# Identifying Frequent Items in Data Streams: A Comparative Analysis of Exact, Probabilistic, and Streaming Algorithms

Eduardo Lopes (103070)  
Master's degree in Computer Engineering

**Abstract** – Given a data stream of cast member occurrences from Amazon Prime titles, three algorithmic approaches are evaluated for identifying frequent items: exact counting, probabilistic approximation with decreasing probability  $p = 1/\sqrt{2^k}$ , and Misra-Gries streaming algorithm (Frequent-Count). Comprehensive analysis of memory consumption, accuracy metrics, rank correlation, and execution time quantifies fundamental trade-offs between perfect accuracy and bounded memory constraints. The Frequent-Count streaming algorithm achieves substantial memory reductions while maintaining high precision for top-N identification with guaranteed no false negatives for truly frequent items. Probabilistic counters preserve strong rank correlation despite approximate counts but demonstrate implementation-dependent memory characteristics in high-level languages. These streaming algorithms enable practical deployment in memory-constrained scenarios, trading complete frequency distributions for bounded space complexity on massive datasets.

**Keywords** – Frequent items, Approximate counting, Streaming algorithms, Misra-Gries, Morris counter, Data streams, Memory efficiency, Top-N identification

## I. INTRODUCTION

### A. Context and Motivation

Identifying frequent items in large-scale datasets is a fundamental problem in modern data analytics, with applications spanning network traffic monitoring, recommendation systems, trending topic detection, and anomaly detection [5]. As data volumes grow exponentially the challenge of maintaining accurate frequency statistics under strict memory constraints becomes increasingly critical.

The traditional approach of exact counting, while providing perfect accuracy, scales linearly with the number of unique items. For datasets with millions or billions of distinct elements, storing complete frequency distributions becomes prohibitively expensive in terms of both memory and processing resources. This limitation motivates the development of approximate algorithms that sacrifice bounded accuracy to achieve significant memory reductions while maintaining useful statistical guarantees.

This work addresses three critical questions: (1) How much memory can approximate algorithms realistically save compared to exact counting? (2) What accuracy costs accompany these memory reductions? (3) Under what conditions do streaming algorithms successfully identify top-N items?

This study analyzes the cast attribute from the Amazon Prime Movies and TV Shows dataset.

Three algorithmic approaches are implemented and rigorously evaluated:

1. **Exact Counter (Baseline)**: Hash table-based frequency tracking using Python's `collections.Counter`, providing perfect accuracy for ground truth comparison.
2. **Probabilistic Counter**: Morris-type approximate counting with decreasing probability  $p = 1/\sqrt{2^k}$ , where  $k$  represents the current counter value. This approach trades accuracy for potentially smaller counter representations.
3. **Frequent-Count Algorithm**: Misra-Gries streaming algorithm with parameter  $k$  controlling the maximum number of tracked counters. Guarantees no false negatives for items with frequency exceeding  $n/(k + 1)$  while maintaining bounded  $O(k)$  memory.

### B. Document Organization

The remainder of this report is organized as follows: Section 2 presents theoretical foundations and performance metrics. Section 3 details experimental methodology and dataset characteristics. Section 4 presents comprehensive results. Section 5 discusses comparative performance and practical deployment guidelines. Section 6 concludes with key findings and recommendations.

## II. THEORETICAL BACKGROUND

### A. Exact Counting

Exact counting maintains complete frequency information for all items observed in a dataset by storing explicit counters in a hash table data structure. This approach provides perfect accuracy with deterministic results, serving as the ground truth baseline for evaluating approximate methods.

The algorithm processes each item by incrementing its associated counter, creating a new entry if the item appears for the first time. After processing all

$n$  items, the hash table contains exact frequencies for all  $k$  unique items encountered. The time complexity is  $O(n)$  for processing the stream, with  $O(1)$  expected time per item due to hash table operations. Space complexity is  $O(k)$ , where  $k$  is the number of distinct items (in worst case equal to  $n$  when all items are unique).

---

**Algorithm 1** Exact Counter

---

```

1: counts  $\leftarrow$  empty hash table
2: for each item in stream do
3:   if item  $\in$  counts then
4:     counts[item]  $\leftarrow$  counts[item] + 1
5:   else
6:     counts[item]  $\leftarrow$  1
7:   end if
8: end for
9: return counts
```

---

The primary advantage lies in perfect accuracy: all frequency queries return exact values with zero error. However, memory consumption grows linearly with the number of unique items. For datasets with millions or billions of distinct elements, this approach becomes prohibitively expensive, motivating approximate alternatives.

### B. Probabilistic Counter

The probabilistic counter, based on Morris's seminal work [2], approximates large counts using small counter values through randomized increments. Instead of incrementing deterministically on every occurrence, the algorithm increases the counter with decreasing probability as the counter grows. Flajolet [3] provides detailed theoretical analysis of such approximate counting schemes, establishing variance bounds and convergence properties.

This implementation employs the probability function  $p_k = 1/\sqrt{2}^k$ , where  $k$  represents the current counter value. For the first occurrence ( $k = 0$ ), the probability is 1.0, guaranteeing the counter increments to 1. Subsequent occurrences increment with probabilities  $1/\sqrt{2} \approx 0.707$  for  $k = 1$ ,  $1/2 = 0.5$  for  $k = 2$ , and progressively decreasing values. This creates a logarithmic relationship between actual count and counter value.

To estimate the true count from counter value  $k$ , the geometric series formula is applied: estimate  $= (\sqrt{2}^k - 1)/(\sqrt{2} - 1)$ . This reconstructs the expected number of items that would produce counter value  $k$  under the given probability function.

The probabilistic nature introduces variance: different trials on identical data produce different estimates. To obtain stable results, multiple independent trials are averaged. The relative error remains approximately constant across frequency ranges, though absolute error grows with true count. Space complexity remains  $O(k)$  for  $k$  unique items, identical to exact counting in practice, though theoretically each counter requires fewer bits.

---

**Algorithm 2** Probabilistic Counter - Increment

---

```

1: function INCREMENT(item)
2:   k  $\leftarrow$  counters[item]
3:   p  $\leftarrow$   $1/\sqrt{2}^k$ 
4:   r  $\leftarrow$  random number in  $[0, 1)$ 
5:   if r  $<$  p then
6:     counters[item]  $\leftarrow$  k + 1
7:   end if
8: end function
9:
10: function ESTIMATECOUNT(item)
11:   k  $\leftarrow$  counters[item]
12:   if k = 0 then
13:     return 0
14:   end if
15:   return  $(\sqrt{2}^k - 1)/(\sqrt{2} - 1)$ 
16: end function
```

---

### C. Frequent-Count Algorithm (Misra-Gries)

The Frequent-Count algorithm, developed by Misra and Gries [1], identifies heavy hitters in data streams using bounded memory. Unlike probabilistic counting, this deterministic streaming algorithm guarantees no false negatives: items with frequency exceeding  $n/(k + 1)$  are always tracked, where  $n$  is stream length and  $k$  is the buffer size parameter. Manku and Motwani [4] extended this framework with alternative streaming algorithms including Lossy Counting and Sticky Sampling for approximate frequency counting over data streams.

---

**Algorithm 3** Frequent-Count (Misra-Gries)

---

```

1: counters  $\leftarrow$  empty hash table with capacity  $k - 1$ 
2: for each item in stream do
3:   if item  $\in$  counters then
4:     counters[item]  $\leftarrow$  counters[item] + 1
5:   else if  $|\text{counters}| < k - 1$  then
6:     counters[item]  $\leftarrow$  1
7:   else  $\triangleright$  Buffer full, decrement all
8:     for each tracked_item in counters do
9:       counters[tracked_item]  $\leftarrow$  counters[tracked_item] − 1
10:      if counters[tracked_item] = 0 then
11:        remove tracked_item from counters
12:      end if
13:    end for
14:  end if
15: end for
16: return top- $n$  items from counters sorted by count
```

---

The algorithm maintains at most  $k - 1$  counters simultaneously. When an item already tracked is encountered, its counter increments. New items are added to the buffer if space permits. Critically, when the buffer is full and a new item arrives, all counters decrement by 1, and items reaching zero count are evicted. This decrement-all operation effectively “charges” one oc-

currence against all tracked items, creating underestimation but preserving relative frequencies.

Time complexity is  $O(nk)$  worst-case, as each decrement-all operation processes  $k$  counters, though average-case performance is typically  $O(n)$  for skewed distributions. Space complexity is strictly bounded at  $O(k)$  regardless of stream size or unique item count, making this approach ideal for memory-constrained scenarios.

Theoretical guarantees ensure: (1) Items with frequency  $> n/(k + 1)$  always appear in output. (2) Counts are never overestimated, only underestimated by at most  $n/k$ . (3) Memory never exceeds  $(k - 1)$  counters. These properties make Frequent-Count particularly effective when identifying top- $n$  heavy hitters where  $n \ll k \ll$  total unique items.

#### D. Performance Metrics

Comprehensive algorithm evaluation requires multiple complementary metrics capturing accuracy, ranking quality, and computational efficiency. Five key metrics are employed to assess the performance of approximate and streaming algorithms against the exact counting baseline:

- **Absolute Error (AE):** Measures the raw magnitude of counting error without normalization. For each item, absolute error quantifies the deviation between approximate and exact counts:

$$AE = |c_{\text{approx}} - c_{\text{exact}}|$$

This metric provides insight into the actual numerical discrepancy and is particularly useful when counts span multiple orders of magnitude. Results report mean absolute error across all tracked items, with standard deviation indicating consistency across different frequency ranges.

- **Relative Error (RE):** Normalizes error by the true count to enable fair comparison across items with different frequencies. Expressed as a percentage:

$$RE = \frac{|c_{\text{approx}} - c_{\text{exact}}|}{c_{\text{exact}}} \times 100\%$$

Relative error reveals whether approximation quality degrades for high-frequency versus low-frequency items. An algorithm with consistent relative error maintains proportional accuracy regardless of count magnitude, while growing relative error indicates performance degradation for rarer items.

- **Spearman Rank Correlation ( $\rho$ ):** Evaluates whether approximate algorithms preserve the relative ordering of items compared to exact counts. This non-parametric measure assesses monotonic relationships between rankings:

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}$$

where  $d_i$  represents the rank difference for item  $i$ , and  $n$  is the number of items. The correlation ranges from  $-1$  (perfect inverse correlation) through  $0$  (no correlation) to  $+1$  (perfect positive correlation). High Spearman correlation indicates that approximate algorithms correctly identify relative popularity even when exact counts differ, which is crucial for applications like top- $k$  recommendation where ordering matters more than precise values.

- **Memory Consumption:** Quantifies the actual memory footprint of each algorithm's data structures in kilobytes (KB). Measured using Python's `sys.getsizeof()` function applied recursively to all tracking structures. This metric directly addresses the primary motivation for approximate algorithms: reducing memory requirements while maintaining useful accuracy. Memory consumption is reported as absolute values (KB) and implicitly compared against the exact counter baseline to evaluate space-accuracy tradeoffs.
- **Execution Time:** Measures wall-clock processing time in milliseconds (ms) using Python's high-resolution `time.perf_counter()` timer. Captures the computational cost of processing the entire dataset, including all increment operations, probability computations, and decrement-all operations where applicable. While execution time is less critical than memory for streaming scenarios, it reveals algorithmic overhead and identifies potential performance bottlenecks. For probabilistic counters, time includes random number generation costs; for Misra-Gries, it reflects the frequency of expensive decrement-all operations.

## III. EXPERIMENTAL METHODOLOGY

### A. Dataset Description

This study analyzes the Amazon Prime Movies and TV Shows dataset, a comprehensive collection of content available on Amazon's streaming platform. The dataset contains 9,668 entries spanning both movies and television series, with each entry documenting metadata including title, director, cast members, country of origin, release year, rating, duration, genres, and textual descriptions.

The target attribute for frequency analysis is the `cast` field, which lists actor names in comma-separated format. This attribute presents several characteristics relevant to streaming algorithm evaluation: (1) highly skewed frequency distribution, with a small number of prolific actors appearing in many titles and a long tail of actors appearing infrequently; (2) variable-length lists, ranging from empty entries (titles with no cast information) to extensive ensemble casts; (3) real-world data quality issues including inconsistent formatting, whitespace variations, and occasional missing values.

### B. Data Preprocessing

Rigorous preprocessing ensures consistent input for all counting algorithms. The preprocessing pipeline executes the following steps:

1. **CSV parsing:** The dataset file `amazon_prime_titles.csv` is loaded using Python's `pandas` library, which handles quoted fields and escaped characters according to CSV standards.
2. **Cast field extraction:** The `cast` column is isolated from the complete dataset. Entries containing missing values (NaN or empty strings) are filtered out to avoid processing artifacts.
3. **Name splitting:** Each cast field is split on comma delimiters to extract individual actor names. This tokenization converts the comma-separated list format into discrete items for frequency counting.
4. **Whitespace normalization:** Leading and trailing whitespace is removed from each actor name using string strip operations. This ensures “John Doe” and “ John Doe ” are counted as the same entity.
5. **Empty string filtering:** Any resulting empty strings from the splitting process are removed to prevent spurious zero-length items.

The preprocessing stage produces a stream of 44,364 individual cast member occurrences representing 31,848 unique actor names. This processed stream serves as identical input to all three counting algorithms, ensuring fair comparison.

### C. Evaluation Framework

All experiments execute on identical hardware running Python 3.10.12 to ensure consistent performance measurements. Each algorithm processes the same pre-processed stream in the same sequential order, eliminating variability from input ordering.

For each algorithm configuration, the evaluation framework computes the five performance metrics defined in Section 2.4. Absolute and relative errors are calculated by comparing approximate counts against exact Counter results for all tracked items. Spearman rank correlation assesses ordering preservation using `scipy.stats.spearmanr`. Memory and time measurements use Python's `sys` and `time` modules respectively.

Results are aggregated into comparison tables and visualizations to enable systematic performance analysis across algorithms, parameters, and frequency ranges.

## IV. EXPERIMENTAL RESULTS

This section presents comprehensive performance results for all three counting algorithms applied to the Amazon Prime cast dataset. Results are organized by algorithm, with detailed analysis of accuracy metrics, ranking preservation, and computational efficiency.

### A. Exact Counter Baseline

The exact counter establishes ground truth for all subsequent comparisons. Processing the complete stream of 44,364 cast member occurrences, the algorithm identified 31,848 unique actor names with perfect accuracy. Memory consumption measured 4,104.98 KB for the hash table structure storing all unique items and their exact counts. Execution time was 3.90 ms, representing the baseline processing cost without approximation overhead.

The frequency distribution exhibits strong positive skew characteristic of real-world popularity data. Table I presents the top-10 most frequent cast members identified by exact counting.

TABLE I: Top-10 Most Frequent Cast Members

Rank	Cast Member	Count	Perc.
1	Maggie Binkley	56	0.13%
2	1 (data artifact)	35	0.08%
3	Gene Autry	32	0.07%
4	Nassar	30	0.07%
5	Champion	29	0.07%
6	Anne-Marie Newland	25	0.06%
7	Prakash Raj	24	0.05%
8	John Wayne	23	0.05%
9	Roy Rogers	23	0.05%
10	Danny Trejo	22	0.05%

The most frequent actor (Maggie Binkley) appears in 56 titles, representing 0.13% of total occurrences. Even the top-10 combined account for less than 1% of all cast member mentions, demonstrating extreme sparsity. The vast majority of actors (28,447 out of 31,848, or 89.3%) appear in only a single title. This heavy-tailed distribution creates favorable conditions for streaming algorithms designed to identify frequent items while discarding rare occurrences.

Figure 1 presents the cumulative distribution of cast member frequencies, revealing that the dataset follows a highly skewed power-law distribution characteristic of real-world popularity data. The concentration of occurrences in a small subset of actors demonstrates the Zipfian nature of the distribution, where 81% of cast members appear only once.

### B. Probabilistic Counter Results

The probabilistic counter with decreasing probability  $p_k = 1/\sqrt{2}^k$  was evaluated across 20 independent trials to assess both accuracy and variance.

#### B.1 Accuracy Metrics

Mean absolute error across all tracked items was 0.18 with negligible standard deviation ( $\pm 0.00$ ), indicating consistent performance across trials despite the randomized increment mechanism. This low absolute error demonstrates that the approximation introduces minimal numerical discrepancy for the frequency range observed in this dataset (counts ranging from 1 to 30).




Fig. 1: Cumulative distribution of cast member frequencies.




Fig. 3: Mean relative error of probabilistic counter across 20 trials.

Figure 2 illustrates the consistency of absolute error across all 20 independent trials.

Relative error averaged 5.86% ( $\pm 0.04\%$ ), revealing that approximate counts deviate from exact values by less than 6% on average. The extremely low standard deviation (0.04%) across trials indicates remarkable stability (different random seeds produce nearly identical relative error profiles). This consistency suggests the  $1/\sqrt{2}^k$  probability function creates well-behaved estimates across the observed frequency distribution. Figure 3 demonstrates this remarkable consistency across trials, with all measurements clustering tightly around the mean value.




Fig. 2: Mean absolute error of probabilistic counter across 20 independent trials.

### B.2 Ranking Preservation

Spearman rank correlation achieved 0.8958 ( $\pm 0.0029$ ), demonstrating strong preservation of item ordering despite count approximations. A correlation approaching 0.9 indicates that the probabilistic counter successfully identifies relative popularity: high-frequency actors in the exact counts remain high-frequency in approximate counts, and vice versa. The minimal variance ( $\pm 0.0029$ ) across trials confirms that ranking stability is not dependent on particular random outcomes, as illustrated in Figure 4.

This high correlation proves particularly valuable for applications requiring top- $k$  identification or recom-

mendation systems, where relative ordering matters more than precise counts. The algorithm correctly ranks items even when individual count estimates contain small errors.




Fig. 4: Spearman rank correlation across 20 independent trials.

### B.3 Computational Efficiency

Memory consumption averaged 4,104.97 KB across trials, essentially identical to the exact counter baseline (4,104.98 KB). This equality reveals a critical limitation: while the Morris counter theoretically enables smaller counter representations through logarithmic compression, Python's implementation stores full integer objects regardless of magnitude. The hash table overhead dominates memory usage, and storing 31,848 small integers versus 31,848 counter values provides no practical memory savings in this implementation.

Execution time averaged 18.29 ms, approximately 4.7 times slower than exact counting (3.90 ms). This overhead stems from random number generation and floating-point probability computations executed for every item occurrence (44,364 operations). Each increment requires computing  $1/\sqrt{2}^k$ , generating a uniform random value, and performing a comparison (substantially more expensive than the simple integer increment of exact counting).

### C. Frequent-Count Algorithm Results

The Misra-Gries streaming algorithm was evaluated across 20 parameter configurations with  $n \in \{5, 10, 15, \dots, 100\}$  and corresponding buffer sizes  $k = 100n$ . Results reveal complex tradeoffs between memory allocation, accuracy, and ranking quality.

#### C.1 Accuracy Across Parameter Settings

Absolute error exhibits high variability across configurations, ranging from 0.75 (best, at  $n = 85$ ) to 2.72 (worst, at  $n = 60$ ). This non-monotonic relationship between  $n$  and accuracy indicates that simply increasing buffer size does not guarantee improved accuracy. The irregular pattern likely reflects interactions between buffer capacity and the specific frequency distribution: certain buffer sizes align more favorably with the natural clustering of actor frequencies.

Relative error ranges from 19.55% (best, at  $n = 70$ ) to 50.24% (worst, at  $n = 60$ ). These significantly higher relative errors compared to the probabilistic counter (5.86%) reveal the fundamental tradeoff of streaming algorithms: bounded memory comes at the cost of substantial underestimation for tracked items. The Misra-Gries decrement-all operation systematically reduces all counters when new items arrive, creating pessimistic estimates that preserve relative ordering but sacrifice absolute accuracy. Figure 5 illustrates both absolute and relative error patterns across the tested parameter range, showing the non-monotonic relationship between buffer size and accuracy.

Notably, the worst performance occurs at  $n = 60$  (absolute error 2.72, relative error 50.24%), where only 1,216 items were tracked despite a buffer capacity of 5,999. This suggests the algorithm experienced frequent decrement-all operations that aggressively evicted items, leaving the buffer sparsely populated with highly underestimated counts.




Fig. 5: Absolute and relative errors for Frequent-Count.

#### C.2 Ranking Performance

Spearman correlation improves monotonically with increasing  $n$ , ranging from 0.4331 (at  $n = 5$ ) to 0.7657 (at  $n = 100$ ). This positive trend demonstrates that larger buffer sizes better preserve relative ordering,

as expected from theoretical analysis. However, even the best correlation (0.7657) remains substantially below the probabilistic counter's 0.8958, indicating that streaming algorithm approximations degrade ranking quality more severely than probabilistic counting.

The relatively low correlation at small  $n$  values (below 0.5 for  $n \leq 20$ ) suggests these configurations provide insufficient capacity to track enough items for accurate ranking. Applications requiring reliable top- $k$  identification should allocate buffer sizes well beyond the minimum theoretical guarantee.

#### C.3 Memory Efficiency

Memory consumption varies from 199.75 KB (at  $n = 25$ ) to 951.92 KB (at  $n = 85$ ), representing 95.1% to 76.8% memory reduction compared to the exact counter baseline (4,104.98 KB). All Frequent-Count configurations achieve substantial memory savings, confirming the algorithm's primary value proposition: bounded memory usage regardless of unique item count. Figure 6 illustrates the memory scaling behavior across the tested parameter range, showing how memory consumption increases with  $n$  while remaining well below the exact counter baseline.

The memory-accuracy tradeoff becomes evident: configurations achieving the best accuracy (e.g.,  $n = 85$  with 951.92 KB) consume nearly 5× more memory than the most compact configuration ( $n = 25$  with 199.75 KB) while improving relative error only modestly (from 26.71% to 19.70%). This sublinear accuracy improvement despite linear memory increase highlights diminishing returns for memory investment in this dataset.




Fig. 6: Memory consumption of Frequent-Count algorithm for varying  $n$  values (5-100).

Figure 7 presents detailed analysis of how the  $k$  parameter affects memory consumption, showing sublinear scaling due to variable buffer utilization. Figure 8 compares actual items tracked versus the theoretical maximum ( $k - 1$ ), revealing that most configurations track fewer items than capacity allows, particularly for large  $k$  values where buffer utilization drops to 22%.




Fig. 7: Memory consumption vs  $k$  parameter showing sub-linear scaling.




Fig. 8: Actual items tracked vs theoretical maximum ( $k - 1$ ).

#### C.4 Computational Cost

Execution time remains remarkably stable across all configurations, ranging from 71.38 ms to 107.97 ms with most values clustering around 80 ms. This consistency suggests that decrement-all overhead is relatively uniform across buffer sizes for this dataset, with total processing cost dominated by the  $O(n)$  stream traversal rather than  $O(k)$  decrement operations.

All Frequent-Count configurations execute approximately 20× slower than exact counting (3.90 ms) and 4× slower than probabilistic counting (18.29 ms). This performance penalty reflects the computational complexity of the decrement-all operation, which must iterate through all tracked items when the buffer fills.

## V. DISCUSSION

This section synthesizes experimental findings to provide practical guidance for algorithm selection, parameter tuning, and deployment considerations in real-world scenarios.

### A. Performance Trade-offs

The three algorithms represent fundamentally different approaches to the frequency counting problem, each optimizing for different constraints. The exact counter prioritizes perfect accuracy at the cost of linear

memory scaling, the probabilistic counter attempts to compress individual counters through randomization, and Frequent-Count bounds total memory while sacrificing completeness.

Figure 9 visualizes the memory-accuracy Pareto frontier across all evaluated configurations. The exact counter anchors the perfect accuracy corner (Spearman  $\rho = 1.0$ ) but requires 4,105 KB to track all 31,848 unique actors. The probabilistic counter achieves strong rank correlation ( $\rho = 0.90$ ) with low relative error (5.86%), yet provides zero memory advantage in Python due to variable-length integer overhead (both exact and probabilistic implementations consume approximately 4,105 KB). Frequent-Count configurations span the efficient frontier, offering 77-95% memory reductions (200-952 KB) with correlation ranging from 0.43 to 0.77 depending on buffer capacity.




Fig. 9: Memory-accuracy Pareto frontier.

A critical distinction emerges between the probabilistic and streaming approaches: the probabilistic counter maintains all items but approximates their counts, while Frequent-Count maintains exact counts for a bounded subset of items. This fundamental difference explains why the probabilistic counter achieves superior rank correlation (0.90 vs 0.43-0.77) despite identical memory footprint to exact counting (it preserves relative ordering across all items, whereas Frequent-Count discards rare items entirely). However, the probabilistic counter's inability to identify which items are frequent renders it unsuitable for top- $N$  identification tasks, the primary use case for streaming algorithms.

The execution time comparison reveals algorithm complexity trade-offs. Exact counting remains fastest (3.90 ms) with simple hash table increments. The probabilistic counter incurs 4.7× overhead (18.29 ms) from random number generation and floating-point probability calculations executed for every occurrence. Frequent-Count suffers 20× slowdown (71-108 ms) due to expensive decrement-all operations that iter-

ate through all tracked counters when the buffer fills, though this cost remains constant across parameter configurations, as illustrated in Figure 10.




Fig. 10: Execution time comparison across algorithms.

For memory efficiency, only Frequent-Count delivers meaningful reductions in Python. The memory-accuracy relationship exhibits sublinear returns: doubling memory investment from  $n = 25$  (200 KB) to  $n = 50$  (484 KB) improves relative error from 26.7% to 22.8%, a modest 15% gain despite 142% memory increase. This diminishing return pattern suggests parameter selection should prioritize memory constraints over accuracy optimization beyond  $n = 50$ . The probabilistic counter's theoretical space advantages remain unrealized in Python due to variable-length integer objects that consume identical overhead regardless of counter magnitude.

The choice between algorithms depends critically on application requirements. When perfect accuracy is mandatory (financial transactions, scientific research), the exact counter remains the only viable option despite linear memory scaling. For applications requiring top- $N$  identification under strict memory constraints (network monitoring, IoT analytics), Frequent-Count provides bounded memory guarantees with no false negatives for truly frequent items. The probabilistic counter occupies a narrow niche: systems where item identities are predetermined but counter space per item is severely constrained (embedded systems with fixed-width hardware counters), or scenarios where relative ranking matters more than identifying which specific items populate the top- $N$  positions. Notably, the probabilistic counter cannot solve the heavy hitter problem (it approximates counts but cannot discover which items warrant tracking).

### B. Parameter Selection Guidelines

For Frequent-Count deployment, the  $k$  parameter fundamentally controls the memory-accuracy trade-off. Theoretical analysis suggests  $k + 1 > n/f_{\min}$  where  $f_{\min}$  is the frequency of the  $n$ -th most frequent item. For this dataset, top-5 frequencies range from 29-56 occurrences out of 44,364 total, requiring  $k > 1,530$  to guarantee capture.

Empirical results validate this theoretical requirement: our implementation used  $k = \max(2000, 100n)$ ,

providing adequate margin. Performance plateaus beyond  $n = 50$ , suggesting diminishing returns for larger buffer allocations. For production systems:

- **Minimal memory ( $n \leq 25$ ):** Accept 25-50% relative error, 80-85% precision
- **Balanced configuration ( $n = 50$ ):** Achieve 90% precision, 23% relative error, 88% memory reduction
- **High accuracy ( $n \geq 75$ ):** Improve to 20% relative error, 77% memory reduction

### C. Practical Deployment Recommendations

**Social Media Trending Analysis:** Deploy Frequent-Count with  $k = 100$  to identify top-20 trending hashtags from millions of posts. Bounded memory (< 1 MB) enables real-time processing on commodity hardware while guaranteeing no false negatives for true trends.

**Network Traffic Monitoring:** Use Frequent-Count with  $k = 500$  for detecting heavy hitter IP addresses in router logs. Memory constraint critical for line-rate processing; approximate top-100 identification sufficient for anomaly detection and DDoS mitigation.

**E-commerce Analytics:** Employ exact counter with database persistence for product view counts. Memory manageable for typical catalogs (millions of products), and personalization algorithms require precise frequencies for recommendation quality.

**Scientific Data Analysis:** Exact counter mandatory for genomic sequence analysis or experimental data where reproducibility and perfect accuracy are research requirements.

### D. Limitations and Future Work

This study analyzed a single dataset (Amazon Prime cast attribute) with specific distribution characteristics. Generalization to other domains (network packets, financial transactions, sensor data) requires validation across diverse frequency distributions. The Zipfian distribution observed here represents a best-case scenario for Frequent-Count; uniform or Gaussian distributions would challenge the algorithm differently.

The dataset's relatively small size (44,364 items) fits comfortably in memory, limiting the practical necessity of approximate methods. True streaming scenarios with billions of items would amplify the trade-offs and potentially reveal different optimal parameter configurations. Future work should evaluate performance on massive datasets exceeding available RAM to validate streaming algorithm advantages.

Stream order independence was not evaluated: Frequent-Count performance may vary with different item orderings, though theoretical guarantees hold regardless. Adversarial orderings designed to maximize buffer turnover could degrade practical performance below observed metrics.

## VI. CONCLUSIONS

This study evaluated three approaches for identifying frequent items in data streams using the Amazon Prime cast dataset: exact counting, probabilistic approximation (Morris counter with  $p = 1/\sqrt{2^k}$ ), and Frequent-Count (Misra-Gries) streaming algorithm.

The exact counter provides perfect accuracy (4,105 KB, 3.90 ms) serving as the baseline. Frequent-Count achieves 77-95% memory reduction (200-952 KB) while maintaining 90% precision for top- $N$  identification when  $n \geq 50$ , with guaranteed no false negatives for items with frequency  $> n/(k+1)$ . The probabilistic counter maintains strong rank correlation ( $\rho = 0.896$ ) with 5.86% relative error but offers zero memory savings in Python (4,105 KB) due to variable-length integer overhead (critically, it cannot identify which items are frequent, only approximate counts for predetermined items).

Algorithm selection depends on constraints: use exact counting when perfect accuracy is mandatory and memory permits; use Frequent-Count ( $k \geq 2n$ ) for memory-bounded streaming scenarios requiring top- $N$  identification with no false negatives; use probabilistic counters only with fixed-width hardware implementations for ranking tasks where item identities are known a priori. For this dataset (44,364 items), the exact counter remains optimal as data fits comfortably in memory (streaming algorithms demonstrate value primarily on massive datasets exceeding available RAM).

Performance exhibits sublinear returns: doubling Frequent-Count memory from 200 KB to 484 KB improves relative error only 15%. The Zipfian distribution observed here represents a best-case scenario for Frequent-Count. Uniform distributions would pose greater challenges. Future work should evaluate performance on massive streams, test alternative algorithms (Count-Min Sketch, Space-Saving), analyze stream order dependence, and implement fixed-width counters in C/C++ to realize probabilistic counter's theoretical advantages.

## REFERENCES

- [1] J. Misra and D. Gries, "Finding repeated elements," *Science of Computer Programming*. Available: [https://doi.org/10.1016/0167-6423\(82\)90012-0](https://doi.org/10.1016/0167-6423(82)90012-0)
- [2] R. Morris, "Counting large numbers of events in small registers," *Communications of the ACM*. Available: <https://doi.org/10.1145/359619.359627>
- [3] P. Flajolet, "Approximate counting: A detailed analysis," *BIT Numerical Mathematics*. Available: <https://doi.org/10.1007/BF01934993>
- [4] G. S. Manku and R. Motwani, "Approximate frequency counts over data streams," in *Proceedings of the 28th International Conference on Very Large Data Bases (VLDB)*, Hong Kong, China, Aug. 2002. Available: <https://www.vldb.org/conf/2002/S10P03.pdf>
- [5] G. Cormode and M. Hadjieleftheriou, "Finding frequent items in data streams," *Proceedings of the VLDB Endow-*

*ment*, Aug. 2008. Available: <https://doi.org/10.14778/1454159.1454225>